

# **The Standard Template Library**



- **The Standard Template Library (STL):** an extensive library of generic templates for classes and functions.
- Categories of Templates:
  - **Containers:** Class templates for objects that store and organize data
  - **Iterators:** Class templates for objects that behave like pointers, and are used to access the individual data elements in a container
  - **Algorithms:** Function templates that perform various operations on elements of containers

## Sequence Containers

- Stores data sequentially in memory, in a fashion similar to an array

## Associative Containers

- Stores data in a non-sequential way that makes it faster to locate elements

**Table 17-1** Sequence Containers

Container Class	Description
array	A fixed-size container that is similar to an array
deque	A double-ended queue. Like a <code>vector</code> , but designed so that values can be quickly added to or removed from the front and back. (This container will be discussed in Chapter 19.)
forward_list	A singly linked list of data elements. Values may be inserted to or removed from any position. (This container will be discussed in Chapter 18.)
list	A doubly linked list of data elements. Values may be inserted to or removed from any position. (This container will be discussed in Chapter 18.)
vector	A container that works like an expandable array. Values may be added to or removed from a <code>vector</code> . The <code>vector</code> automatically adjusts its size to accommodate the number of elements it contains.

## Containers

**Table 17-2** Associative Containers

Container Class	Description
<code>set</code>	Stores a set of unique values that are sorted. No duplicates are allowed.
<code>multiset</code>	Stores a set of unique values that are sorted. Duplicates are allowed.
<code>map</code>	Maps a set of keys to data elements. Only one key per data element is allowed. Duplicates are not allowed. The elements are sorted in order of their keys.
<code>multimap</code>	Maps a set of keys to data elements. Many keys per data element are allowed. Duplicates are allowed. The elements are sorted in order of their keys.
<code>unordered_set</code>	Like a <code>set</code> , except that the elements are not sorted
<code>unordered_multiset</code>	Like a <code>multiset</code> , except that the elements are not sorted
<code>unordered_map</code>	Like a <code>map</code> , except that the elements are not sorted
<code>unordered_multimap</code>	Like a <code>multimap</code> , except that the elements are not sorted



## Container Adapters

**Table 17-3** Container Adapter Classes

Container Adapter Class	Description
stack	An adapter class that stores elements in a deque (by default). A stack is a last-in, first-out (LIFO) container. When you retrieve an element from a stack, the stack always gives you the last element that was inserted. (This class will be discussed in Chapter 19.)
queue	An adapter class that stores elements in a deque (by default). A queue is a first-in, first-out (FIFO) container. When you retrieve an element from a stack, the stack always gives you the first, or earliest, element that was inserted. (This class will be discussed in Chapter 19.)
priority_queue	An adapter class that stores elements in a vector (by default). A data structure in which the element that you retrieve is always the element with the greatest value. (This class will be discussed in Chapter 19.)



**Table 17-4** Header Files

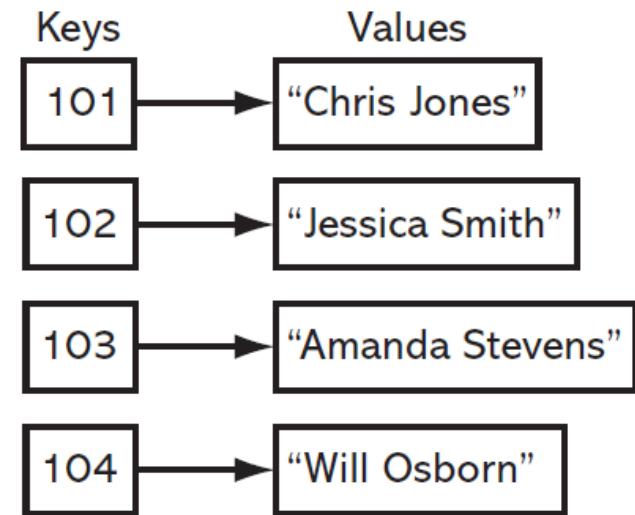
Header File	Classes
<array>	array
<deque>	deque
<forward_list>	forward_list
<list>	list
<map>	map, multimap
<queue>	queue, priority_queue
<set>	set, multiset
<stack>	stack
<unordered_map>	unordered_map, unordered_multimap
<unordered_set>	unordered_set, unordered_multiset
<vector>	vector

# **The map, multimap, and unordered\_map Classes**



## Maps general concept

- A *map* is an associative container
- Each element that is stored in a map has two parts: a *key* and a *value*.
- To retrieve a specific value from a map, you use the key that is associated with that value.
- This is similar to the process of looking up a word in the dictionary, where the words are keys and the definitions are values.
- Example: a map in which employee IDs are the keys and employee names are the values.
- You use an employee's ID to look up that employee's name.



- You can use the STL map class to store key-value pairs.
- The keys that are stored in a map container are unique – no duplicates.
- The map class is declared in the `<map>` header file.

### Default Constructor

```
map<keyDataType, valueType> name;  
Creates an empty map.
```

### Range Constructor

```
map<keyDataType, valueType>  
name(iterator1, iterator2);  
Creates a map that is initialized with a range of values  
from another map. iterator1 marks the beginning of the  
range and iterator2 marks the end.
```

### Copy Constructor

```
map<keyDataType, valueType> name(map2);  
Creates a map that is a copy of map2.
```

## The map class

- Example: defining a map container to hold employee ID numbers (as ints) and their corresponding employee names (as strings):

```
map<int, string> employees;
```

Key data type      Value data type

```
map<int, string> employees =  
{  
    {101, "Chris Jones"}, {102, "Jessica Smith"},  
    {103, "Amanda Stevens"}, {104, "Will Osborn"}  
};
```

- In the first element, the key is 101 and the value is "Chris Jones".
- In the second element, the key is 102 and the value is "Jessica Smith".
- In the third element, the key is 103 and the value is "Amanda Stevens".
- In the fourth element, the key is 104 and the value is "Will Osborn".



## The overloaded [ ] operator

- You can use the [ ] operator to add new elements to a map.
- General format:

*mapName[key] = value;*

- This adds the key-value pair to the map.
- If the key already exists in the map, its associated value will be changed to *value*.

```
map<int, string> employees;
employees[110] = "Beth Young";
employees[111] = "Jake Brown";
employees[112] = "Emily Davis";
```

- After this code executes, the `employees` map will contain the following elements:
  - Key = 110, Value = "Beth Young"
  - Key = 111, Value = "Jake Brown"
  - Key = 112, Value = "Emily Davis"



## The pair type

- Internally, the elements of a `map` are stored as instances of the `pair` type.
- `pair` is a struct that has two member variables: `first` and `second`.
- The element's key is stored in `first`, and the element's value is stored in `second`.
- The `pair` struct is declared in the `<utility>` header file. When you `#include` the `<map>` header file, `<utility>` is automatically included as well.



## Inserting an element using the insert function

- The `map` class provides an `insert()` member function that adds a `pair` object as an element to the `map`.
- You can use the STL function template `make_pair` to construct a `pair` object.
- The `make_pair` function template is declared in the `<utility>` header file.

```
map<int, string> employees;
employees.insert(make_pair(110, "Beth Young"));
employees.insert(make_pair(111, "Jake Brown"));
employees.insert(make_pair(112, "Emily Davis"));
```

- After this code executes, the `employees` map will contain the following elements:
  - Key = 110, Value = "Beth Young"
  - Key = 111, Value = "Jake Brown"
  - Key = 112, Value = "Emily Davis"

**Note:** If the element that you are inserting with the `insert()` member function has the same key as an existing element, the function will *not* insert the new element.



## retrieving an element from a map

- You can use the `at()` member function to retrieve a map element by its key:

```
// Create a map containing employee IDs and names.  
map<int, string> employees = {  
    {101, "Chris Jones"}, {102, "Jessica Smith"},  
    {103, "Amanda Stevens"}, {104, "Will Osborn"}  
};  
// Retrieve a value from the map.  
cout << employees.at(103) << endl;
```

Displays "Amanda Stevens"

- To prevent the `at()` member function from throwing an exception (if the specified key does not exist), use the `count` member function to determine whether it exists:

```
// Create a map containing employee IDs and names.  
map<int, string> employees = {  
    {101, "Chris Jones"}, {102, "Jessica Smith"},  
    {103, "Amanda Stevens"}, {104, "Will Osborn"}  
};  
  
// Retrieve a value from the map.  
if (employees.count(103))  
    cout << employees.at(103) << endl;  
else  
    cout << "Employee not found.\n";
```

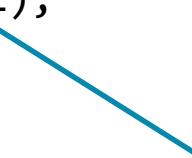
The `count()` member function returns 1 if the specified key exists, or 0 otherwise.



## Deleting an element from a map

- You can use the `erase()` member function to retrieve a map element by its key:

```
// Create a map containing employee IDs and names.  
map<int, string> employees =  
{  
    {101, "Chris Jones"}, {102, "Jessica Smith"},  
    {103, "Amanda Stevens"}, {104, "Will Osborn"}  
};  
  
// Delete the employee with ID 102.  
employees.erase(102);
```



Deletes Jessica Smith from the map



## Range loops and maps

```
// Create a map containing employee IDs and names.  
map<int, string> employees =  
{  
    {101, "Chris Jones"}, {102, "Jessica Smith"},  
    {103, "Amanda Stevens"}, {104, "Will Osborn"}  
};
```

**Remember, each element is a pair.**

```
// Display each element.  
for (pair<int, string> element : employees)  
{  
    cout << "ID: " << element.first << "\tName: "  
        << element.second << endl;  
}
```

```
for (auto element : employees) ←————— auto simplifies this  
{  
    cout << "ID: " << element.first << "\tName: "  
        << element.second << endl;  
}
```



- The `begin()` and `end()` member functions return a bidirectional iterator of the `iterator` type
- The `cbegin()` and `cend()` member functions return a bidirectional iterator of the `const_iterator` type
- The `rbegin()` and `rend()` member functions return a reverse bidirectional iterator of the `reverse_iterator` type
- The `crbegin()` and `crend()` member functions return a reverse bidirectional iterator of the `const_reverse_iterator` type

## Iterators and maps

- When an iterator points to a map element, it points to an instance of the pair type.
- The element has two member variables: first and second.
- The element's key is stored in first, and the element's value is stored in second.

### Program 17-11

```
1 // This program demonstrates an iterator with a map.
2 #include <iostream>
3 #include <string>
4 #include <map>
5 using namespace std;
6
7 int main()
8 {
9     // Create a map containing employee IDs and names.
10    map<int, string> employees =
11        { {101,"Chris Jones"}, {102,"Jessica Smith"}, 
12          {103,"Amanda Stevens"},{104,"Will Osborn"} };
13
14    // Create an iterator.
15    map<int, string>::iterator iter;
16
17    // Use the iterator to display each element in the map.
18    for (iter = employees.begin(); iter != employees.end(); iter++)
19    {
20        cout << "ID: " << iter->first
21                      << "\tName: " << iter->second << endl;
22    }
23
24    return 0;
25 }
```

### Program Output

```
ID: 101 Name: Chris Jones
ID: 102 Name: Jessica Smith
ID: 103 Name: Amanda Stevens
ID: 104 Name: Will Osborn
```

## Storing Objects Of Your Own Classes as *Values* in a map

- If you want to store an object as a value in a map, there is one requirement for that object's class:

● It must have a default constructor.

- Consider the following Contact class...

```
1 #ifndef CONTACT_H
2 #define CONTACT_H
3 #include <string>
4 using namespace std;
5
6 class Contact
7 {
8 private:
9     string name;
10    string email;
11 public:
12     Contact() ← Default constructor
13     {   name = "";
14         email = ""; }
15
16     Contact(string n, string em)
17     {   name = n;
18         email = em; }
19
20     void setName(string n)
21     {   name = n; }
22
23     void setEmail(string em)
24     {   email = em; }
25
26     string getName() const
27     {   return name; }
28
29     string getEmail() const
30     {   return email; }
31 };
32 #endif
```



## Program 17-14

```
1 #include <iostream>
2 #include <string>
3 #include <map>
4 #include "Contact.h"
5 using namespace std;
6
7 int main()
8 {
9     string searchName;    // The name to search for
10
11    // Create some Contact objects
12    Contact contact1("Ashley Miller", "amiller@faber.edu");
13    Contact contact2("Jacob Brown", "jbrown@gotham.edu");
14    Contact contact3("Emily Ramirez", "eramirez@coolidge.edu");
15
16    // Create a map to hold the Contact objects.
17    map<string, Contact> contacts;
18
19    // Create an iterator for the map.
20    map<string, Contact>::iterator iter;
21
22    // Add the contact objects to the map.
23    contacts[contact1.getName()] = contact1;
24    contacts[contact2.getName()] = contact2;
25    contacts[contact3.getName()] = contact3;
```

In the map, the keys are the contact names, and the values are the Contact objects.



## Storing Objects Of Your Own Classes as *Values* in a map

```
26
27     // Get the name to search for.
28     cout << "Enter a name: ";
29     getline(cin, searchName);
30
31     // Search for the name.
32     iter = contacts.find(searchName);
33
34     // Display the results.
35     if (iter != contacts.end())
36     {
37         cout << "Name: " << iter->second.getName() << endl;
38         cout << "Email: " << iter->second.getEmail() << endl;
39     }
40     else
41     {
42         cout << "Contact not found.\n";
43     }
44
45     return 0;
46 }
```

### Program Output (with Example Input Shown in Bold)

Enter a name: **Emily Ramirez**

Name: Emily Ramirez

Email: eramirez@coolidge.edu

### Program Output (with Example Input Shown in Bold)

Enter a name: **Billy Clark**

Contact not found.



## Storing Objects Of Your Own Classes as key in a map

- If you want to store an object as a key in a map, there is one requirement for that object's class:

**It must overload the < operator.**

- Consider the following Customer class...

```
1 #ifndef CUSTOMER_H
2 #define CUSTOMER_H
3 #include<string>
4 using namespace std;
5
6 class Customer
7 {
8 private:
9     int custNumber;
10    string name;
11 public:
12     Customer(int cn, string n)
13     { custNumber = cn;
14         name = n; }
15
16     void setCustNumber(int cn)
17     { custNumber = cn; }
18
19     void setName(string n)
20     { name = n; }
21
22     int getCustNumber() const
23     { return custNumber; }
24
25     string getName() const
26     { return name; }
27
28     bool operator < (const Customer &right) const
29     { bool status = false;
30
31         if (custNumber < right.custNumber)
32             status = true;
33
34         return status; }
35     };
36 #endif
```



# Storing Objects Of Your Own Classes as key in a map

## Program 17-17

```
1 #include <iostream>
2 #include <string>
3 #include <map>
4 #include "Customer.h"
5 using namespace std;
6
7 int main()
8 {
9     // Create some Customer objects.
10    Customer customer1(1001, "Sarah Scott");
11    Customer customer2(1002, "Austin Hill");
12    Customer customer3(1003, "Megan Cruz");
13
14    // Create a map to hold the seat assignments.
15    map<Customer, string> assignments;
16
17    // Use the map to store the seat assignments.
18    assignments[customer1] = "1A";
19    assignments[customer2] = "2B";
20    assignments[customer3] = "3C";
21
22    // Display all objects in the map.
23    for (auto element : assignments)
24    {
25        cout << element.first.getName() << "\t"
26            << element.second << endl;
27    }
28
29    return 0;
30 }
```

## Program Output

Sarah Scott	1A
Austin Hill	2B
Megan Cruz	3C

This program assigns seats in a theater to customers. The map uses Customer objects as keys, and seat numbers as values.



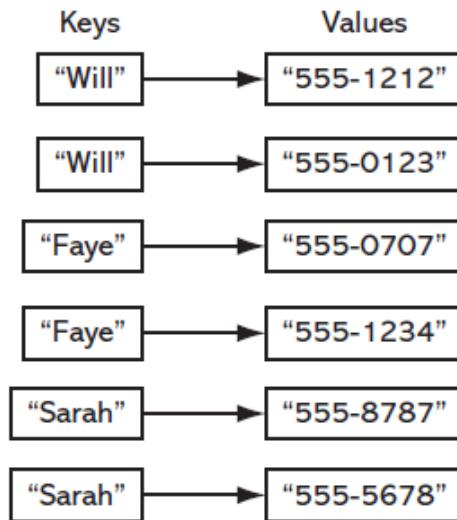
## Unordered\_map class

- The `unordered_map` class is similar to the `map` class, except in two regards:
  - The keys in an `unordered_map` are not sorted
  - The `unordered_map` class has better performance
- You should use the `unordered_map` class instead of the `map` class if:
  - You will be making a lot of searches on a large number of elements
  - You are not concerned with retrieving them in key order
- The `unordered_map` class is declared in the `<unordered_map>` header file



## multimap class

- The `multimap` class is a map that allows duplicate keys
- The `multimap` class has most of the same member functions as the `map` class (see Table 17-11 in your textbook)
- The `multimap` class is declared in the `<map>` header file
- Consider a phonebook application where the key is a person's name and the value is that person's phone number.
- A multimap container would allow each person to have multiple phone numbers



# multimap class

## Program 17-19

```
1 #include <iostream>
2 #include <string>
3 #include <map>
4 using namespace std;
5
6 int main()
7 {
8     // Define a phonebook multimap.
9     multimap<string, string> phonebook =
10        { {"Will", "555-1212"}, {"Will", "555-0123"} ,
11         {"Faye", "555-0707"}, {"Faye", "555-1234"} ,
12         {"Sarah", "555-8787"}, {"Sarah", "555-5678"} };
13
14     // Display the elements in the multimap.
15     for (auto element : phonebook)
16     {
17         cout << element.first << "\t"
18             << element.second << endl;
19     }
20     return 0;
21 }
```

## Program Output

Faye 555-0707  
Faye 555-1234  
Sarah 555-8787  
Sarah 555-5678  
Will 555-1212  
Will 555-0123



## Adding an element to a multimap

- The `multimap` class does not overload the `[]` operator.
  - So, you cannot use an assignment statement to add a new element to a `multimap`.
- Instead, you will use either the `emplace()` or the `insert()` member functions.

```
multimap<string, string> phonebook;
phonebook.emplace("Will", "555-1212");
phonebook.emplace("Will", "555-0123");
phonebook.emplace("Faye", "555-0707");
phonebook.emplace("Faye", "555-1234");
phonebook.emplace("Sarah", "555-8787");
phonebook.emplace("Sarah", "555-5678");
```



## Getting the number of elements with the same key

- The multimap class's count() member function accepts a key as its argument, and returns the number of elements that match the specified key.

### Program 17-20

```
1 #include <iostream>
2 #include <string>
3 #include <map>
4 using namespace std;
5
6 int main()
7 {
8     // Define a phonebook multimap.
9     multimap<string, string> phonebook =
10        { {"Will", "555-1212"}, {"Will", "555-0123"} ,
11         {"Faye", "555-0707"}, {"Faye", "555-1234"} ,
12         {"Sarah", "555-8787"}, {"Sarah", "555-5678"} };
13
14     // Display the number of elements that match "Faye".
15     cout << "Faye has " << phonebook.count("Faye") << " elements.\n";
16     return 0;
17 }
```

### Program Output

Faye has 2 elements.



## Retrieving an element with specific key

- The multimap class has a `find()` member function that searches for an element with a specified key.
- The `find()` function returns an iterator to the first element matching it.
- If the element is not found, the `find()` function returns an iterator to the end of the multimap.
- To retrieve all elements matching a specified key, use the `equal_range` member function.
- The `equal_range` member function returns a pair object.
- The pair object's first member is an iterator pointing to the first element that matches the specified key.
- The pair object's second member is an iterator pointing to the position after the last element that matches the specified key.



## Retrieving an element with specific key

```
// Define a phonebook multimap.  
multimap<string, string> phonebook =  
    { {"Will", "555-1212"}, {"Will", "555-0123"},  
      {"Faye", "555-0707"}, {"Faye", "555-1234"},  
      {"Sarah", "555-8787"}, {"Sarah", "555-5678"} };  
  
// Define a pair variable to receive the object that  
// is returned from the equal_range member function.  
pair<multimap<string, string>::iterator,  
     multimap<string, string>::iterator> range;  
  
// Define an iterator for the multimap.  
multimap<string, string>::iterator iter;  
  
// Get the range of elements that match "Faye".  
range = phonebook.equal_range("Faye");  
  
// Display all of the elements that match "Faye".  
for (iter = range.first; iter != range.second; iter++)  
{  
    cout << iter->first << "\t" << iter->second << endl;  
}
```

### Program Output

```
Faye      555-0707  
Faye      555-1234
```



## Deleting an element with specific key

- To delete all elements matching a specified key, use the `erase()` member function.

```
// Define a phonebook multimap.  
multimap<string, string> phonebook =  
    { {"Will", "555-1212"}, {"Will", "555-0123"},  
     {"Faye", "555-0707"}, {"Faye", "555-1234"},  
     {"Sarah", "555-8787"}, {"Sarah", "555-5678"} };  
  
// Delete Will's phone numbers from the multimap.  
phonebook.erase("Will");
```



- The `unordered_multimap` class is similar to the `multimap` class, except:
  - The keys in an `unordered_multimap` are not sorted
  - The `unordered_multimap` class has better performance
- You should use the `unordered_multimap` class instead of the `multimap` class if:
  - You will be making a lot of searches on a large number of elements
  - You are not concerned with retrieving them in key order
- The `unordered_multimap` class is declared in the `<unordered_multimap>` header file

# **The `set`, `multiset`, and `unordered_set` Classes**



- A set is an associative container that is similar to a mathematical set.
- You can use the STL set class to create a set container.
- All the elements in a set must be unique. No two elements can have the same value.
- The elements in a set are automatically sorted in ascending order.
- The set class is declared in the `<set>` header file.

## The set class

- You can use the STL set class to create a set container.
- The keys that are stored in a map container are unique – no duplicates.
- The map class is declared in the `<map>` header file.

**Default  
Constructor**      `set<dataType> name;`  
**Creates an empty set.**

Range Constructor      `set<dataType> name(iterator1, iterator2);`  
**Creates a set that is initialized with a range of values.**  
***iterator1* marks the beginning of the range and**  
***iterator2* marks the end.**

Copy Constructor      `set<dataType> name(set2);`  
**Creates a set that is a copy of *set2*.**



- Example: defining a set container to hold integers:

```
set<int> numbers;
```

- Example: defining and initializing a set container to hold integers:

```
set<int> numbers = {1, 2, 3, 4, 5};
```

- A set cannot contain duplicate items.
- If the same value appears more than once in an initialization list, it will be added to the set only one time.
- For example, the following set will contain the values 1, 2, 3, 4, and 5:

```
set<int> numbers = {1, 1, 2, 2, 3, 4, 5, 5, 5};
```



- The insert() member function adds a new element to a set:

```
set<int> numbers;
numbers.insert(10);
numbers.insert(20);
numbers.insert(30);
```

```
// Create a set containing names.
set<string> names = {"Joe", "Karen", "Lisa", "Jackie"};

// Display each element.
for (string element : names)
{
    cout << element << endl;
```

- The `begin()` and `end()` member functions return a bidirectional iterator of the `iterator` type
- The `cbegin()` and `cend()` member functions return a bidirectional iterator of the `const_iterator` type
- The `rbegin()` and `rend()` member functions return a reverse bidirectional iterator of the `reverse_iterator` type
- The `crbegin()` and `crend()` member functions return a reverse bidirectional iterator of the `const_reverse_iterator` type

```
// Create a set containing names.  
set<string> names = {"Joe", "Karen", "Lisa", "Jackie"};  
  
// Create an iterator.  
set<string>::iterator iter;  
  
// Use the iterator to display each element in the set.  
for (iter = names.begin(); iter != names.end(); iter++)  
{  
    cout << *iter << endl;  
}
```

## Determine whether an element exists

- The set class's count() member function accepts a value as its argument, and returns 1 if that value exists in the set. The function returns 0 otherwise.

```
set<string> names = {"Joe", "Karen", "Lisa", "Jackie"};
if (names.count("Lisa"))
    cout << "Lisa was found in the set.\n";
else
    cout << "Lisa was not found.\n";
```



## Retrieving an element

- The set class has a `find()` member function that searches for an element with a specified value.
- The `find()` function returns an iterator to the element matching it.
- If the element is not found, the `find()` function returns an iterator to the end of the set.

```
// Create a set containing names.  
set<string> names = {"Joe", "Karen", "Lisa", "Jackie"};  
  
// Create an iterator.  
set<string>::iterator iter;  
  
// Find "Karen".  
iter = names.find("Karen");  
  
// Display the result.  
if (iter != names.end())  
{  
    cout << *iter << " was found.\n";  
}  
else  
{  
    cout << "Karen was not found.\n";  
}
```



- If you want to store an object in a set, there is one requirement for that object's class:

It must overload the < operator.

- Consider the following Customer class...

```
25     string getName() const
26     { return name; }
27
28     bool operator < (const Customer &right) const
29     { bool status = false;
30
31         if (custNumber < right.custNumber)
32             status = true;
33
34         return status; }
35     };
36 #endif
```

```
1     #ifndef CUSTOMER_H
2     #define CUSTOMER_H
3     #include<string>
4     using namespace std;
5
6     class Customer
7     {
8     private:
9         int custNumber;
10        string name;
11    public:
12        Customer(int cn, string n)
13        { custNumber = cn;
14            name = n; }
15
16        void setCustNumber(int cn)
17        { custNumber = cn; }
18
19        void setName(string n)
20        { name = n; }
21
22        int getCustNumber() const
23        { return custNumber; }
24
```

# Storing objects of your own classes in a set

## Program 17-22

```
1 #include <iostream>
2 #include <set>
3 #include "Customer.h"
4 using namespace std;
5
6 int main()
7 {
8     // Create a set of Customer objects.
9     set<Customer> customerset =
10        { Customer(1003, "Megan Cruz"),
11          Customer(1002, "Austin Hill"),
12          Customer(1001, "Sarah Scott")
13      };
14
15     // Try to insert a duplicate customer number.
16     customerset.emplace(1001, "Evan Smith");
17
18     // Display the set elements
19     cout << "List of customers:\n";
20     for (auto element : customerset)
21     {
22         cout << element.getCustNumber() << " "
23             << element.getName() << endl;
24     }
25 }
```

```
26     // Search for customer number 1002.
27     cout << "\nSearching for Customer Number 1002:\n";
28     auto it = customerset.find(Customer(1002, ""));
29
30     if (it != customerset.end())
31         cout << "Found: " << it->getName() << endl;
32     else
33         cout << "Not found.\n";
34
35     return 0;
36 }
```

## Program Output

List of customers:

1001 Sarah Scott

1002 Austin Hill

1003 Megan Cruz

Searching for Customer Number 1002:

Found: Austin Hill



- The `multiset` class is a set that allows duplicate items.
- The `multiset` class has the same member functions as the `set` class (see Table 17-13 in your textbook).
- The `multiset` class is declared in the `<set>` header file.
- In the `set` class, the `count()` member function returns either 0 or 1. In the `multiset` class, the `count()` member function can return values greater than 1.
- In the `set` class, the `equal_range()` member function returns a range with, at most, one element. In the `multiset` class, the `equal_range()` member function can return a range with multiple elements.

## Unordered\_set class

- The unordered\_set class is similar to the set class, except in two regards:
  - The values in an unordered\_set are not sorted
  - The unordered\_set class has better performance
- You should use the unordered\_set class instead of the set class if:
  - You will be making a lot of searches on a large number of elements
  - You are not concerned with retrieving them in ascending order
- The unordered\_set class is declared in the <unordered\_set> header file



# Algorithms



# STL Algorithms

- The STL provides a number of algorithms, implemented as function templates, in the `<algorithm>` header file.
- These functions perform various operations on ranges of elements.
- A range of elements is a sequence of elements denoted by two iterators:
  - The first iterator points to the first element in the range
  - The second iterator points to the end of the range (the element to which the second iterator points is not included in the range).



# Categories of Algorithms in the STL

- Min/max algorithms
- Sorting algorithms
- Search algorithms
- Read-only sequence algorithms
- Copying and moving algorithms
- Swapping algorithms
- Replacement algorithms
- Removal algorithms
- Reversal algorithms
- Fill algorithms
- Rotation algorithms
- Shuffling algorithms
- Set algorithms
- Transformation algorithm
- Partition algorithms
- Merge algorithms
- Permutation algorithms
- Heap algorithms
- Lexicographical comparison algorithm



# Sorting

- ➊ The sort function:

```
sort(iterator1, iterator2);
```

*iterator1* and *iterator2* mark the beginning and end of a range of elements. The function sorts the range of elements in ascending order.



# Searching

- The `binary_search` function:

```
binary_search(iterator1, iterator2, value);
```

*iterator1* and *iterator2* mark the beginning and end of a range of elements that are sorted in ascending order. *value* is the value to search for. The function returns true if *value* is found in the range, or false otherwise.



## Program 17-23

```
1 #include <iostream>
2 #include <vector>
3 #include <algorithm>
4 using namespace std;
5
6 int main()
7 {
8     int searchValue; // Value to search for
9
10    // Create a vector of unsorted integers.
11    vector<int> numbers = {10, 1, 9, 2, 8, 3, 7, 4, 6, 5};
12
13    // Sort the vector.
14    sort(numbers.begin(), numbers.end());
15}
```

```
16    // Display the vector.
17    cout << "Here are the sorted values:\n";
18    for (auto element : numbers)
19        cout << element << " ";
20    cout << endl;
21
22    // Get the value to search for.
23    cout << "Enter a value to search for: ";
24    cin >> searchValue;
25
26    // Search for the value.
27    if (binary_search(numbers.begin(), numbers.end(), searchValue))
28        cout << "That value is in the vector.\n";
29    else
30        cout << "That value is not in the vector.\n";
31
32    return 0;
33 }
```

### Program Output

Here are the sorted values:

1 2 3 4 5 6 7 8 9 10

Enter a value to search for: 8

That value is in the vector.

### Program Output

Here are the sorted values:

1 2 3 4 5 6 7 8 9 10

Enter a value to search for: 99

That value is not in the vector.



# Detecting Permutations

- If a range has  $N$  elements, there are  $N!$  possible arrangements, or permutations, of those elements.
- For example, the range of integers 1, 2, 3 has six possible permutations:

1, 2, 3  
1, 3, 2  
2, 1, 3  
2, 3, 1  
3, 1, 2  
3, 2, 1



# Detecting Permutations

- The `is_permutation()` function determines whether one range of elements is a permutation of another range of elements.

`is_permutation(iterator1, iterator2, iterator3)`

- `iterator1` and `iterator2` mark the beginning and end of the first range of elements.
- `iterator3` marks the beginning of the second range of elements, assumed to have the same number of elements as the first range.
- The function returns true if the second range is a permutation of the first range, or false otherwise.
- See Program 17-25 in your textbook for an example.



# Plugging Your Own Functions into an Algorithm

- Many of the function templates in the STL are designed to accept function pointers as arguments.
- This allows you to “plug” one of your own functions into the algorithm.
- For example:

```
for_each(iterator1, iterator2, function)
```

- iterator1* and *iterator2* mark the beginning and end of a range of elements.
- function* is the name of a function that accepts an element as its argument.
- The `for_each()` function iterates over the range of elements, passing each element as an argument to *function*.



# Plugging Your Own Functions into an Algorithm

```
void doubleNumber(int &n)
{
    n = n * 2;
}
```

```
vector<int> numbers = { 1, 2, 3, 4, 5 };

// Display the numbers before doubling.
for (auto element : numbers)
    cout << element << " ";
cout << endl;

// Double the value of each vector element.
for_each(numbers.begin(), numbers.end(), doubleNumber);

// Display the numbers before doubling.
for (auto element : numbers)
    cout << element << " ";
cout << endl;
```



# Plugging Your Own Functions into an Algorithm

- Another example:

```
count_if(iterator1, iterator2, function)
```

- *iterator1* and *iterator2* mark the beginning and end of a range of elements.
- *function* is the name of a function that accepts an element as its argument, and returns either true or false.
- The `count_if()` function iterates over the range of elements, passing each element as an argument to *function*.
- The `count_if` function returns the number of elements for which *function* returns true.



```
1 #include <iostream>
2 #include <vector>
3 #include <algorithm>
4 using namespace std;
5
6 // Function prototypes
7 bool isNegative(int);
8
9 int main()
10 {
11     // Create a vector of ints.
12     vector<int> numbers = { 0, 99, 120, -33, 10, 8, -1, 101 };
13
14     // Get the number of elements that are negative.
15     int negatives = count_if(numbers.begin(), numbers.end(), isNegative);
16
17     // Display the results.
18     cout << "There are " << negatives << " negative elements.\n";
19     return 0;
20 }
21
22 // isNegative function
23 bool isNegative(int n)
24 {
25     bool status = false;
26
27     if (n < 0)
28         status = true;
29
30     return status;
31 }
```

### **Program Output**

There are 2 negative elements.



# Algorithms for Set Operations

- The STL provides function templates for basic mathematical set operations.

STL Function Template	Description
set_union	Finds the union of two sets, which is a set that contains all the elements of both sets, excluding duplicates.
set_intersection	Finds the intersection of two sets, which is a set that contains only the elements that are found in both sets.
set_difference	Finds the difference of two sets, which is the set of elements that appear in one set, but not the other.
set_symmetric_difference	Finds the symmetric difference of two sets, which is the set of elements that appear in one set, but not both.
set_includes	Determines whether one set includes another.



# **Introduction to Function Objects and Lambda Expressions**



# Function Objects

- A function object is an object that acts like a function.
  - It can be called
  - It can accept arguments
  - It can return a value
- Function objects are also known as *functors*



# Function Objects

- Orange To create a function object, you write a class that overloads the () operator.

```
1 #ifndef SUM_H
2 #define SUM_H
3
4 class Sum
5 {
6 public:
7     int operator()(int a, int b)
8     { return a + b; } ← Returns an int
9 };
10#endif
```

Accepts two `int` arguments



## Program 17-34

```
1 #include <iostream>
2 #include "Sum.h"
3 using namespace std;
4
5 int main()
6 {
7     // Local variables
8     int x = 10;
9     int y = 2;
10    int z = 0;
11
12    // Create a Sum object.
13    Sum sum;
14
15    // Call the sum function object.
16    z = sum(x, y);
17
18    // Display the result.
19    cout << z << endl;
20
21    return 0;
22 }
```

## Program Output

12



# Anonymous Function Objects

- Function objects can be called at the point of their creation, without being given a name. Consider this class:

```
1 #ifndef IS EVEN_H
2 #define IS EVEN_H
3
4 class IsEven
5 {
6 public:
7     bool operator()(int x)
8     { return x % 2 == 0; }
9 };
10#endif
```



## Program 17-36

```
1 #include <iostream>
2 #include <vector>
3 #include <algorithm>
4 #include "IsEven.h"
5 using namespace std;
6
7 int main()
8 {
9     // Create a vector of ints.
10    vector<int> v = { 1, 2, 3, 4, 5, 6, 7, 8 };
11
12    // Get the number of elements that even.
13    int evenNums = count_if(v.begin(), v.end(), IsEven());
14
15    // Display the results.
16    cout << "The vector contains " << evenNums << " even numbers.\n";
17
18 }
```

An `IsEven` object is created here, but not given a name.  
It is anonymous.

## Program Output

The vector contains 4 even numbers.



# Predicate Terminology

- A function or function object that returns a Boolean value is called a *predicate*.
- A predicate that takes only one argument is called a *unary predicate*.
- A predicate that takes two arguments is called a *binary predicate*.
- This terminology is used in much of the available C++ documentation and literature.



# Lambda Expressions

- A lambda expression is a compact way of creating a function object without having to write a class declaration.
- It is an expression that contains only the logic of the object's operator() member function.
- When the compiler encounters a lambda expression, it automatically generates a function object in memory, using the code that you provide in the lambda expression for the operator() member function.



# Lambda Expressions

- General format:

`[](parameter list) { function body }`

- The [ ] is known as the lambda introducer. It marks the beginning of a lambda expression.
- parameter list* is a list of parameter declarations for the function object's operator() member function.
- function body* is the code that should be the body of the object's operator() member function.



# Lambda Expressions

- Example: a lambda expression for a function object that computes the sum of two integers:

```
[](int a, int b) { return a + b; }
```

- Example: a lambda expression for a function object that determines whether an integer is even is:

```
[](int x) { return x % 2 == 0; }
```

- Example: a lambda expression for a function object that takes an integer as input and prints the square of that integer:

```
[](int a) { cout << a * a << " "; }
```



# Lambda Expressions

- When you call a lambda expression, you write a list of arguments, enclosed in parentheses, right after the expression.
- For example, the following code snippet displays 7, which is the sum of the variables x and y:

```
int x = 2;  
int y = 5;  
cout << [](int a, int b) {return a + b;}(x, y) << endl;
```



# Lambda Expressions

- The following code segment counts the even numbers in a vector:

```
// Create a vector of ints.  
vector<int> v = { 1, 2, 3, 4, 5, 6, 7, 8 };  
  
// Get the number of elements that are even.  
int evenNums = count_if(v.begin(), v.end(), [](int x) {return x % 2 == 0;});  
  
// Display the results.  
cout << "The vector contains " << evenNums << " even numbers.\n";
```



# Lambda Expressions

- Because lambda expressions generate function objects, you can assign a lambda expression to a variable and then call it through the variable's name:

```
auto sum = [](int a, int b) {return a + b;};  
int x = 2;  
int y = 5;  
int z = sum(x, y);
```



## Program 17-37

```
1 #include <iostream>
2 #include <vector>
3 #include <algorithm>
4 using namespace std;
5
6 int main()
7 {
8     // Create a vector of ints.
9     vector<int> v = { 1, 2, 3, 4, 5, 6, 7, 8 };
10
11    // Use a lambda expression to create a function object.
12    auto isEven = [] (int x) { return x % 2 == 0; };
13
14    // Get the number of elements that even.
15    int evenNums = count_if(v.begin(), v.end(), isEven);
16
17    // Display the results.
18    cout << "The vector contains " << evenNums << " even numbers.\n";
19    return 0;
20 }
```

## Program Output

The vector contains 4 even numbers.



# Functional Classes in the STL

- The STL library defines a number of classes that you can instantiate to create function objects in your program.
- To use these classes, you must `#include` the `<functional>` header file.
- Table 17-15 in your textbook lists a few of the functional classes:

**Table 17-15** STL Function Object Classes

Functional Class	Description
<code>less&lt;T&gt;</code>	<code>less&lt;T&gt;()</code> ( <code>T a, T b</code> ) is true if and only if $a < b$
<code>less_equal&lt;T&gt;</code>	<code>less_equal()</code> ( <code>T a, T b</code> ) is true if and only if $a \leq b$
<code>greater&lt;T&gt;</code>	<code>greater&lt;T&gt;()</code> ( <code>T a, T b</code> ) is true if and only if $a > b$
<code>greater_equal&lt;T&gt;</code>	<code>greater_equal&lt;T&gt;()</code> ( <code>T a, T b</code> ) is true if and only if $a \geq b$

